

Predicting Quarterback 2nd Contract Value

Paul Sabin

Contents

0.1	1. Introduction	1
0.2	2. Data Loading and Preprocessing	1
0.3	3. Modeling Approach	2
0.4	4. Model Specification and Training	3
0.5	5. Model Evaluation	4
0.6	6. Variable Importance	6
0.7	7. Summary	6

0.1 1. Introduction

This analysis builds a **random forest regression model** to predict a quarterback's **second contract value (APY as a percentage of salary cap)**.

College and NFL performance metrics, physical traits, and draft data are used as predictors.

We use **k-fold cross-validation** to estimate performance and produce: - Variable importance plots

- Cross-validated metrics (RMSE, MAE, Bias)
- Visualization of predicted vs actual performance
- Review of largest over- and under-predictions

0.2 2. Data Loading and Preprocessing

The entire script is in this file:

```
source("predict_qb_contract.R")
```

```
## read in cfb qb usage data from 2013
## read in cfb qb usage data from 2014
## read in cfb qb usage data from 2015
## read in cfb qb usage data from 2016
## read in cfb qb usage data from 2017
## read in cfb qb usage data from 2018
## read in cfb qb usage data from 2019
## read in cfb qb usage data from 2020
## read in cfb qb usage data from 2021
## read in cfb qb usage data from 2022
## read in cfb qb usage data from 2023
## read in cfb qb usage data from 2024
```

To confirm the structure:

```
data_to_model %>%  
  glimpse()
```

```
## Rows: 111  
## Columns: 33  
## $ gsis_id      <chr> "00-0030998", "00-0031064", "00-0031076", "00-00312~  
## $ player       <chr> "Keith Wenning", "Tom Savage", "David Fales", "Tedd~  
## $ college      <chr> "Ball State", "Pittsburgh", "San Jose State", "Loui~  
## $ conference   <chr> "MAC", "ACC", "MWC", "AAC", "ACC", "MWC", "SEC", "A~  
## $ birth_date   <date> 1991-02-14, 1990-04-26, 1990-10-04, 1992-11-10, 19~  
## $ draft_year   <int> 2014, 2014, 2014, 2014, 2014, 2014, 2014, 2014, 201~  
## $ draft_number <dbl> 194, 135, 183, 32, 120, 36, 163, 214, 3, 1, 147, 98~  
## $ undrafted    <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~  
## $ apy_cap_pct  <dbl> 0.003, 0.008, 0.004, 0.034, 0.004, 0.150, 0.001, 0.~  
## $ height       <dbl> 74, 76, 73, 74, 78, 75, 73, 76, 77, 76, 75, 73, 73,~  
## $ weight       <dbl> 223, 230, 219, 215, 250, 215, 207, 221, 245, 231, 2~  
## $ rookie_age   <dbl> 24.54483, 26.35181, 24.90897, 22.80630, 24.16975, 2~  
## $ usg_overall  <dbl> 0.543, 0.497, 0.503, 0.519, 0.544, 0.591, 0.472, 0.~  
## $ usg_pass     <dbl> 0.973, 0.945, 0.907, 0.973, 0.872, 0.920, 0.896, 0.~  
## $ usg_rush     <dbl> 0.067, 0.077, 0.063, 0.096, 0.255, 0.068, 0.086, 0.~  
## $ usg_1st_down <dbl> 0.504, 0.390, 0.441, 0.399, 0.464, 0.557, 0.431, 0.~  
## $ usg_2nd_down <dbl> 0.519, 0.517, 0.553, 0.568, 0.544, 0.608, 0.457, 0.~  
## $ usg_3rd_down <dbl> 0.768, 0.748, 0.595, 0.801, 0.745, 0.694, 0.697, 0.~  
## $ usg_standard_downs <dbl> 0.518, 0.414, 0.434, 0.432, 0.478, 0.536, 0.419, 0.~  
## $ usg_passing_downs <dbl> 0.609, 0.673, 0.662, 0.728, 0.672, 0.740, 0.599, 0.~  
## $ passing_att  <dbl> 454, 376, 487, 382, 391, 605, 347, 504, 351, 422, 3~  
## $ passing_cmp  <dbl> 296, 230, 312, 268, 224, 424, 225, 335, 239, 276, 2~  
## $ passing_pct  <dbl> 0.652, 0.612, 0.641, 0.702, 0.573, 0.701, 0.648, 0.~  
## $ passing_yds  <dbl> 3933, 2834, 4189, 3523, 2861, 4866, 3075, 3528, 328~  
## $ passing_td   <dbl> 34, 21, 33, 28, 16, 48, 26, 21, 22, 24, 21, 12, 30,~  
## $ passing_int  <dbl> 6, 9, 13, 4, 13, 7, 9, 7, 7, 17, 5, 5, 16, 6, 11, 2~  
## $ passing_ypa  <dbl> 8.7, 7.5, 8.6, 9.2, 7.3, 8.0, 8.9, 7.0, 9.3, 8.4, 8~  
## $ rushing_att  <dbl> 40, 72, 48, 57, 159, 40, 53, 83, 79, 49, 148, 83, 8~  
## $ rushing_yds  <dbl> 45, -208, 7, 54, 295, 117, 186, 267, 179, 80, 548, ~  
## $ rushing_td   <dbl> 5, 3, 2, 0, 4, 2, 7, 6, 5, 3, 8, 4, 2, 0, 5, 14, 4,~  
## $ rushing_ypc  <dbl> 1.1, -2.9, 0.1, 0.9, 1.9, 2.9, 3.5, 3.2, 2.3, 1.6, ~  
## $ rushing_long <dbl> 11, 12, 16, 20, 26, 17, 57, 53, 20, 28, 29, 21, 35,~  
## $ rush_pct     <dbl> 0.08097166, 0.16071429, 0.08971963, 0.12984055, 0.2~
```

0.3 3. Modeling Approach

We model `apy_cap_pct` using all variables **from conference onward** as predictors. Before modeling the log-transformation is taken then converted back before evaluation.

A **random forest** is used because it handles nonlinearities, interactions, and mixed variable types naturally.

0.4 4. Model Specification and Training

```
# Define the recipe
rf_recipe <- recipe(apv_cap_pct ~ ., data = data_to_model) |>
  update_role(gsis_id, player, college, birth_date, draft_number, new_role = "ID") |>
  step_rm(gsis_id, player, college, birth_date) |>
  step_dummy(all_nominal_predictors()) |>
  step_zv(all_predictors()) |>
  step_normalize(all_numeric_predictors())

# Random forest model spec
rf_spec <- rand_forest(mtry = tune(), min_n = tune(), trees = 500) |>
  set_mode("regression") |>
  set_engine("ranger", importance = "permutation")

# Workflow
rf_workflow <- workflow() |>
  add_model(rf_spec) |>
  add_recipe(rf_recipe)

# Cross-validation folds
folds <- vfold_cv(data_to_model, v = 5)

# Tune model
rf_tune <- tune_grid(
  rf_workflow,
  resamples = folds,
  grid = 10,
  metrics = metric_set(rmse, mae, rsq)
)

# Select best parameters
best_params <- select_best(rf_tune, metric = "rmse")

# Finalize workflow
final_rf_workflow <- finalize_workflow(rf_workflow, best_params)

# === Fit resamples to get out-of-fold predictions ===
cv_fit <- fit_resamples(
  final_rf_workflow,
  resamples = folds,
  control = control_resamples(save_pred = TRUE)
)

# --- Collect out-of-fold predictions ---
cv_preds <- collect_predictions(cv_fit)
```

```

# Join with original data (if needed for player names)
cv_preds <- cv_preds |>
  left_join(data_to_model |>
    dplyr::mutate(rowid = 1:n()) |>
    dplyr::select(rowid, player), by = c(".row" = "rowid"))

# === Evaluation Metrics ===
cv_metrics <- cv_preds |> metrics(truth = apy_cap_pct, estimate = .pred)
cv_bias <- cv_preds |> summarise(bias = mean(.pred - apy_cap_pct))

```

0.5 5. Model Evaluation

0.5.1 5.1 Cross-Validated Metrics

```

cv_metrics %>%
  dplyr::select(-.estimator) |>
  bind_rows(
    tibble(.estimate = cv_bias$bias, .metric = "bias")
  ) %>%
  kable(digits = 4, caption = "Cross-Validated Performance Metrics") %>%
  kable_styling(full_width = FALSE)

```

Table 1: Cross-Validated Performance Metrics

.metric	.estimate
rmse	0.0753
rsq	0.0203
mae	0.0360
bias	-0.0277

0.5.2 5.2 Predicted vs Actual Plot

```

cv_preds %>%
  ggplot(aes(y = apy_cap_pct, x = .pred)) +
  geom_point(alpha = 0.7) +
  geom_smooth(method = 'glm', color = 'red', se = FALSE) +
  geom_abline(slope = 1, intercept = 0, color = "black") +
  xlim(0, 0.25) + ylim(0, 0.25) +
  labs(
    title = "Predicted vs Actual Contract Value",
    x = "Predicted APY (as % of Cap)",
    y = "Actual APY (as % of Cap)"
  ) +
  theme(aspect.ratio = 1)

```

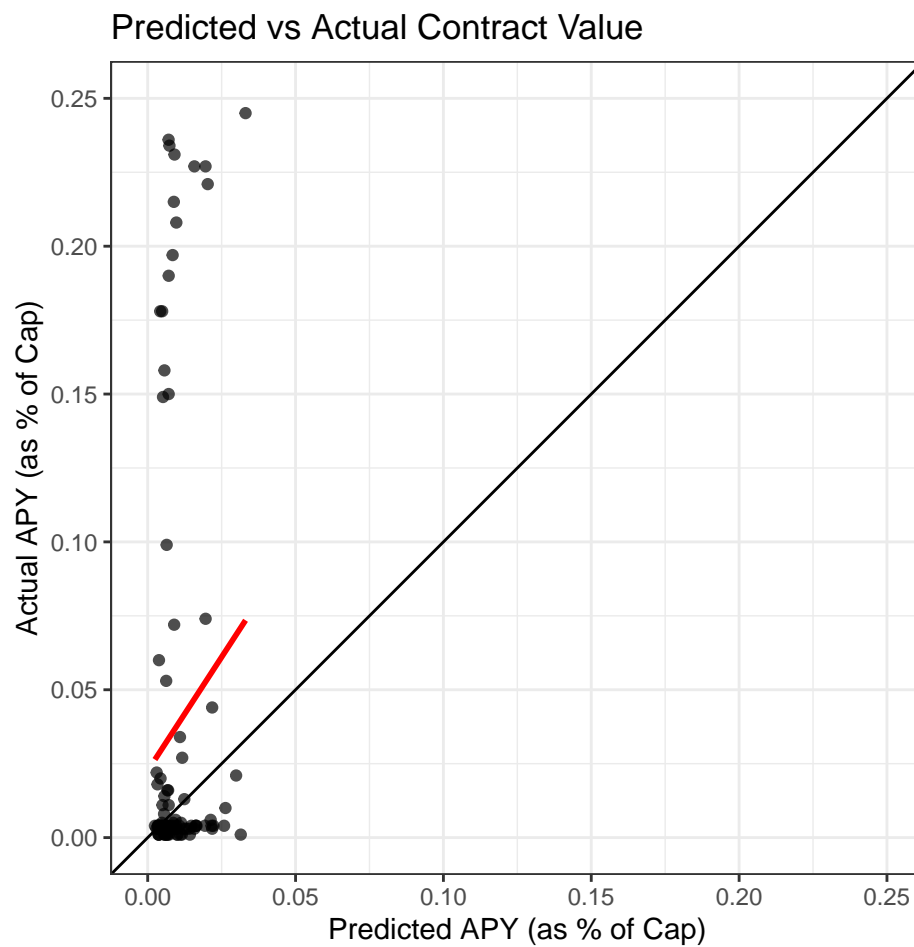


Figure 1: Predicted vs Actual apy_cap_pct

0.5.3 5.3 Biggest Over/Under Predictions

```
cv_preds %>%  
  mutate(error = .pred - apy_cap_pct, abs_error = abs(error)) %>%  
  arrange(desc(abs_error)) %>%  
  select(player, apy_cap_pct, .pred, error) %>%  
  head(10) %>%  
  kable(digits = 4, caption = "Largest Model Misses (Over/Under Predictions)") %>%  
  kable_styling(full_width = FALSE)
```

Table 2: Largest Model Misses (Over/Under Predictions)

player	apy_cap_pct	.pred	error
Josh Allen	0.236	0.0071	-0.2289
Justin Herbert	0.234	0.0074	-0.2266
Lamar Jackson	0.231	0.0091	-0.2219
Joe Burrow	0.245	0.0331	-0.2119
Patrick Mahomes	0.227	0.0158	-0.2112
Jalen Hurts	0.227	0.0196	-0.2074
Trevor Lawrence	0.215	0.0088	-0.2062
Kyler Murray	0.221	0.0203	-0.2007
Tua Tagovailoa	0.208	0.0097	-0.1983
Deshaun Watson	0.197	0.0084	-0.1886

0.6 6. Variable Importance

```
final_rf_fit <- fit(final_rf_workflow, data = data_to_model)  
final_rf_fit %>%  
  extract_fit_parsnip() %>%  
  vip(num_features = 20)
```

0.7 7. Summary

This report used a random forest to predict quarterback second-contract values as a proportion of the salary cap.

Key takeaways: - The model achieves weak predictive accuracy (see RMSE and MAE above). - Key predictors often include age (selection bias), passing numbers, and running efficiency. - The largest outliers are quarterbacks who ended up being among the best in the NFL because they get paid so much better than anyone else.

```
## R version 4.4.2 (2024-10-31)  
## Platform: aarch64-apple-darwin20  
## Running under: macOS Sequoia 15.7.1  
##  
## Matrix products: default  
## BLAS: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
```




Figure 2: Top 20 Most Important Features

```

## LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## time zone: America/New_York
## tzcode source: internal
##
## attached base packages:
## [1] tools      splines  stats      graphics  grDevices  utils      datasets
## [8] methods    base
##
## other attached packages:
## [1] ranger_0.17.0      tidybayes_3.0.7      broom.mixed_0.2.9.6
## [4] mice_3.17.0        cfbfastR_1.9.5        nflfastR_5.1.0
## [7] magrittr_2.0.3      lme4_1.1-37           Matrix_1.7-3
## [10] xfun_0.52          kableExtra_1.4.0      vip_0.4.1
## [13] yardstick_1.3.2     workflowsets_1.1.1    workflows_1.2.0
## [16] tune_1.3.0          rsample_1.3.1         recipes_1.3.0
## [19] parsnip_1.3.2       modeldata_1.4.0       infer_1.0.9
## [22] dials_1.4.1         scales_1.4.0          broom_1.0.8
## [25] tidymodels_1.3.0    lubridate_1.9.4       forcats_1.0.0
## [28] stringr_1.5.1       dplyr_1.1.4           purrr_1.0.4
## [31] readr_2.1.5         tidyr_1.3.1           tibble_3.2.1
## [34] ggplot2_4.0.0       tidyverse_2.0.0
##
## loaded via a namespace (and not attached):
## [1] RColorBrewer_1.1-3  tensorA_0.36.2.1      rstudioapi_0.17.1
## [4] jsonlite_2.0.0      shape_1.4.6.1         jomo_2.7-6
## [7] farver_2.1.2        nloptr_2.2.1          rmarkdown_2.29
## [10] fs_1.6.6            vctrs_0.6.5           memoise_2.0.1
## [13] minqa_1.2.8         tinytex_0.57          janitor_2.2.1
## [16] sparsevctrs_0.3.3   htmltools_0.5.8.1     curl_6.2.3
## [19] distributional_0.5.0 xgboost_1.7.11.1      mitml_0.4-5
## [22] parallelly_1.44.0   cachem_1.1.0          lifecycle_1.0.4
## [25] iterators_1.0.14     pkgconfig_2.0.3       R6_2.6.1
## [28] fastmap_1.2.0        rbibutils_2.3         future_1.49.0
## [31] snakecase_0.11.1     digest_0.6.37         furrr_0.3.1
## [34] textshaping_1.0.1    labeling_0.4.3        timechange_0.3.0
## [37] abind_1.4-8          httr_1.4.7            mgcv_1.9-3
## [40] compiler_4.4.2       withr_3.0.2           S7_0.2.0
## [43] backports_1.5.0      pan_1.9               MASS_7.3-65
## [46] lava_1.8.1           future.apply_1.11.3    nnet_7.3-20
## [49] glue_1.8.0          nlme_3.1-168          grid_4.4.2
## [52] checkmate_2.3.2      generics_0.1.4        gtable_0.3.6
## [55] tzdb_0.5.0          class_7.3-23          data.table_1.17.2
## [58] hms_1.1.3           xml2_1.3.8            foreach_1.5.2
## [61] pillar_1.10.2        ggdist_3.3.3          nflreadr_1.5.0

```


## [64]	posterior_1.6.1	lhs_1.2.0	lattice_0.22-7
## [67]	sfd_0.1.0	survival_3.8-3	tidyselect_1.2.1
## [70]	knitr_1.50	reformulas_0.4.1	arrayhelpers_1.1-0
## [73]	fastrmodels_2.0.0	svglite_2.2.1	hardhat_1.4.1
## [76]	timeDate_4041.110	stringi_1.8.7	DiceDesign_1.10
## [79]	yaml_2.3.10	boot_1.3-31	evaluate_1.0.3
## [82]	codetools_0.2-20	cli_3.6.5	RcppParallel_5.1.10
## [85]	rpart_4.1.24	systemfonts_1.2.3	Rdpack_2.6.4
## [88]	Rcpp_1.0.14	globals_0.18.0	coda_0.19-4.1
## [91]	svUnit_1.0.6	parallel_4.4.2	gower_1.0.2
## [94]	GPfit_1.0-9	listenv_0.9.1	glmnet_4.1-8
## [97]	viridisLite_0.4.2	ipred_0.9-15	prodlim_2025.04.28
## [100]	rlang_1.1.6		